




Normalization

Why Normalize?




$$o = w x$$

$$\frac{\partial \ell(o)}{\partial w} = \left(\frac{d \ell(o)}{d o} \right) x^T$$


Scalar
derivative

$$x_i > 0$$

$$\left(\frac{\partial \ell(o)}{\partial w} \right)_i \geq 0 \quad \left(\frac{\partial \ell(o)}{\partial w} \right)_i \leq 0$$



Why Normalize?




$$o = w x$$

$$\frac{\partial \ell(o)}{\partial w} = \left(\frac{d \ell(o)}{d o} \right) x^T$$

$$x \in \mathbb{R}^2 \quad |x_1| \ll |x_2|$$

$$\frac{\partial \ell(o)}{\partial w_1} \ll \frac{\partial \ell(o)}{w_2}$$



Input Normalization

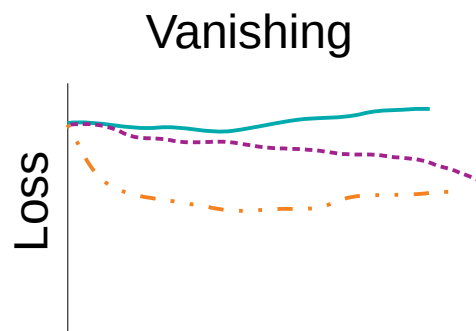

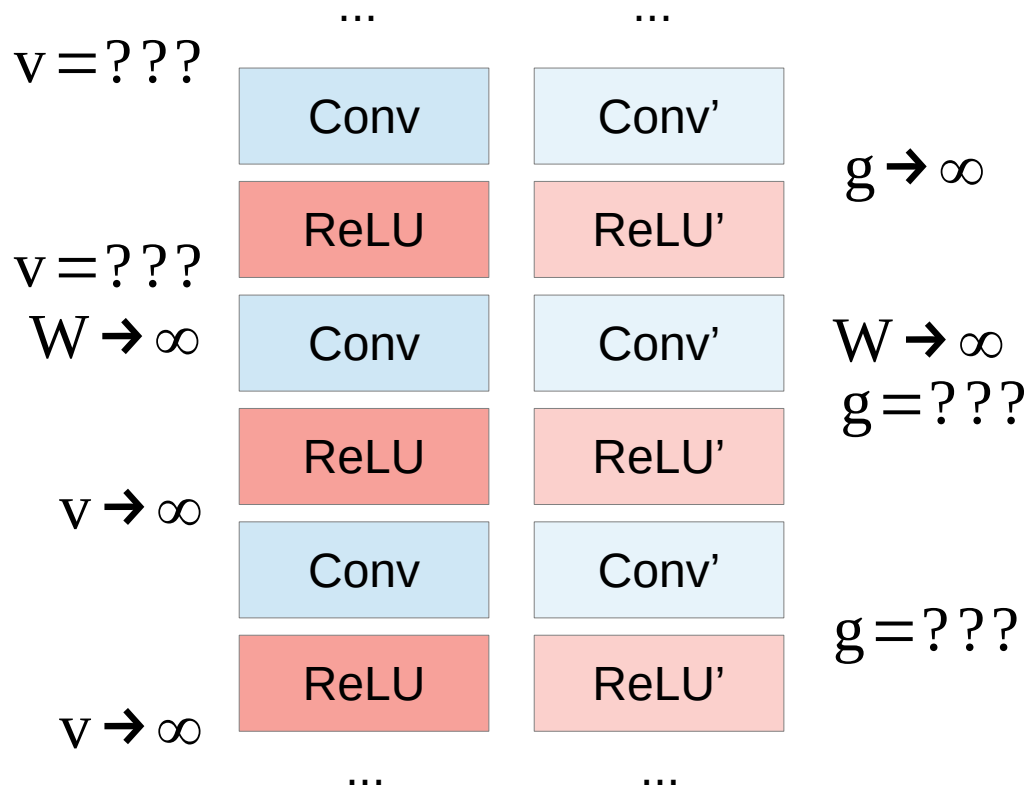
 x_i

$$\frac{x_i - \mu_x}{\sigma_x}$$

Average over
all elements of
all inputs


For images, compute mean and standard deviation for each channel – that is, one red mean, one blue mean, and one green mean.

Vanishing / Exploding Gradients



$$\|g_i\| \ll \|W_i\|$$

Normalization



$$y = \alpha x + \beta$$

$$E[y] = 0$$

$$\text{Var}[y] = 1$$

Batch Normalization

Conv

BatchNorm

ReLU

$$y = \alpha x + \beta$$

Over the entire
batch

$$\mathbb{E}[y] = 0$$

$$\text{Var}[y] = 1$$

$$\mathbf{x} \in \mathbb{R}^{B \times C \times H \times W}$$

$$y_{i,c,j,k} = \frac{x_{i,c,j,k} - \mu_c}{\sigma_c}$$

$$\mu_c = \frac{1}{BHW} \sum_{i,j,k} x_{i,c,j,k}$$

$$\sigma_c^2 = \frac{1}{BHW} \sum_{i,j,k} (x_{i,c,j,k} - \mu_c)^2$$

Batch Normalization

- ✓ Keeps the activation magnitudes in check
- ✓ Deals with badly scaled weights
- ✗ Mixes gradient information between inputs
 - Mitigated by large batches

$$\mathbf{X} \in \mathbb{R}^{B \times C \times H \times W}$$

$$X_{i,c,j,k} \rightarrow \infty$$

$$\mu_c \rightarrow \infty \quad \sigma_c \rightarrow \infty$$

BatchNorm at Test Time

- Usually we don't test on a batch of data.
- Keep a running average of the mean and standard deviation during training, then save those values.

Layer Normalization

- Same as BatchNorm, but we compute statistics per input rather than per channel.
- Prevents cross-talk.
- Training and testing are the same.
- In practice, works well for sequence models but not in computer vision.
 - We need separate channels.

$$\mathbf{x} \in \mathbb{R}^{B \times C \times H \times W}$$

$$\mu_i = \frac{1}{CHW} \sum_{c,j,k} x_{i,c,j,k}$$

$$\sigma_i^2 = \frac{1}{CHW} \sum_{c,j,k} (x_{i,c,j,k} - \mu_i)^2$$

Instance Normalization

- Compute statistics per input *and* per channel
 - Sum over *only* spacial locations
- Works okay for image generating in computer graphics
- Not so good in recognition
- Statistics are unstable

$$\mathbf{x} \in \mathbb{R}^{B \times C \times H \times W}$$

$$\mu_{i,c} = \frac{1}{HW} \sum_{j,k} x_{i,c,j,k}$$

$$\sigma_{i,c}^2 = \frac{1}{HW} \sum_{j,k} (x_{i,c,j,k} - \mu_{i,c})^2$$

Group Normalization

- Compute statistics over groups of channels
 - Between instance normalization and layer normalization
- More flexible than layer normalization, more stable than instance normalization.

$$\mathbf{X} \in \mathbb{R}^{B \times C \times H \times W}$$


$$S = \lfloor C / G \rfloor$$

$$\mu_{i,g} = \frac{1}{SHW} \sum_{j,k} \sum_{c=S(g-1)}^{Sg-1} \mathbf{X}_{i,c,j,k}$$


$$\sigma_{i,g}^2 = \frac{1}{SHW} \sum_{j,k} \sum_{c=S(g-1)}^{Sg-1} \left(\mathbf{X}_{i,c,j,k} - \mu_{i,g} \right)^2$$

Summary


Batch normalization




Instance normalization




Layer normalization



Group normalization



Normalization in Practice



- No Bias needed in Conv
- Activations are zero mean
 - ReLU will zero out half of activations
- Learn a scale and bias parameter in the normalization layer (`affine=True`)

- Scale and bias in the normalization layer are optional.
- Conv is unchanged

NOTE: Do not normalize after linear layers (statistical estimates are too unstable)